

Introduction to Deep Learning

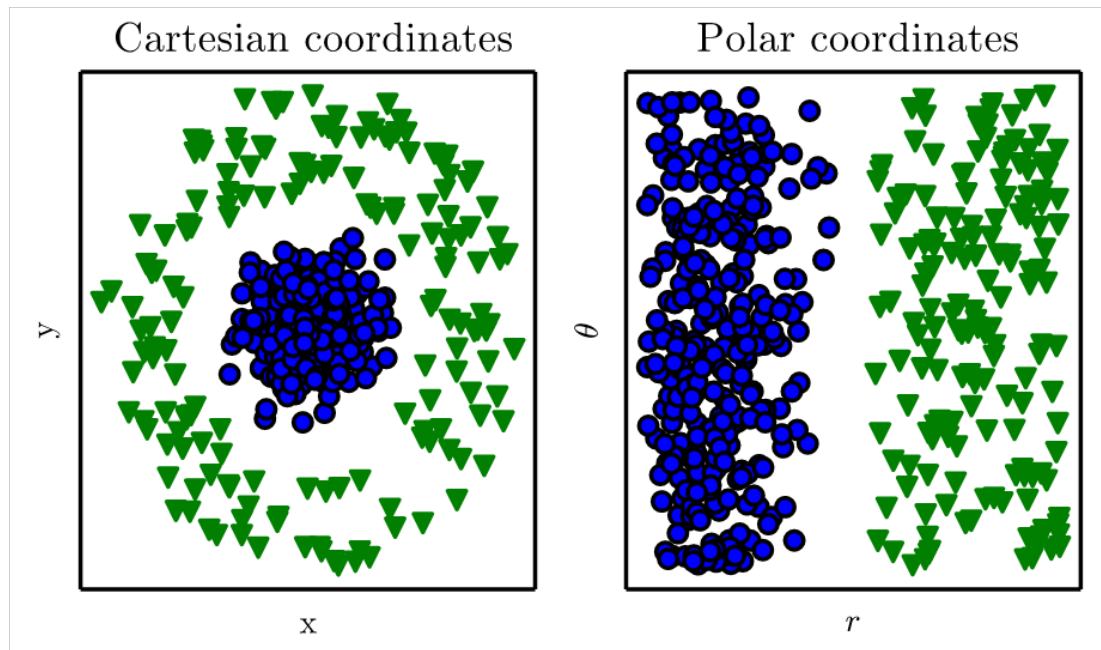
CMPT 733

Steven Bergner
sbergner@sfu.ca

Overview

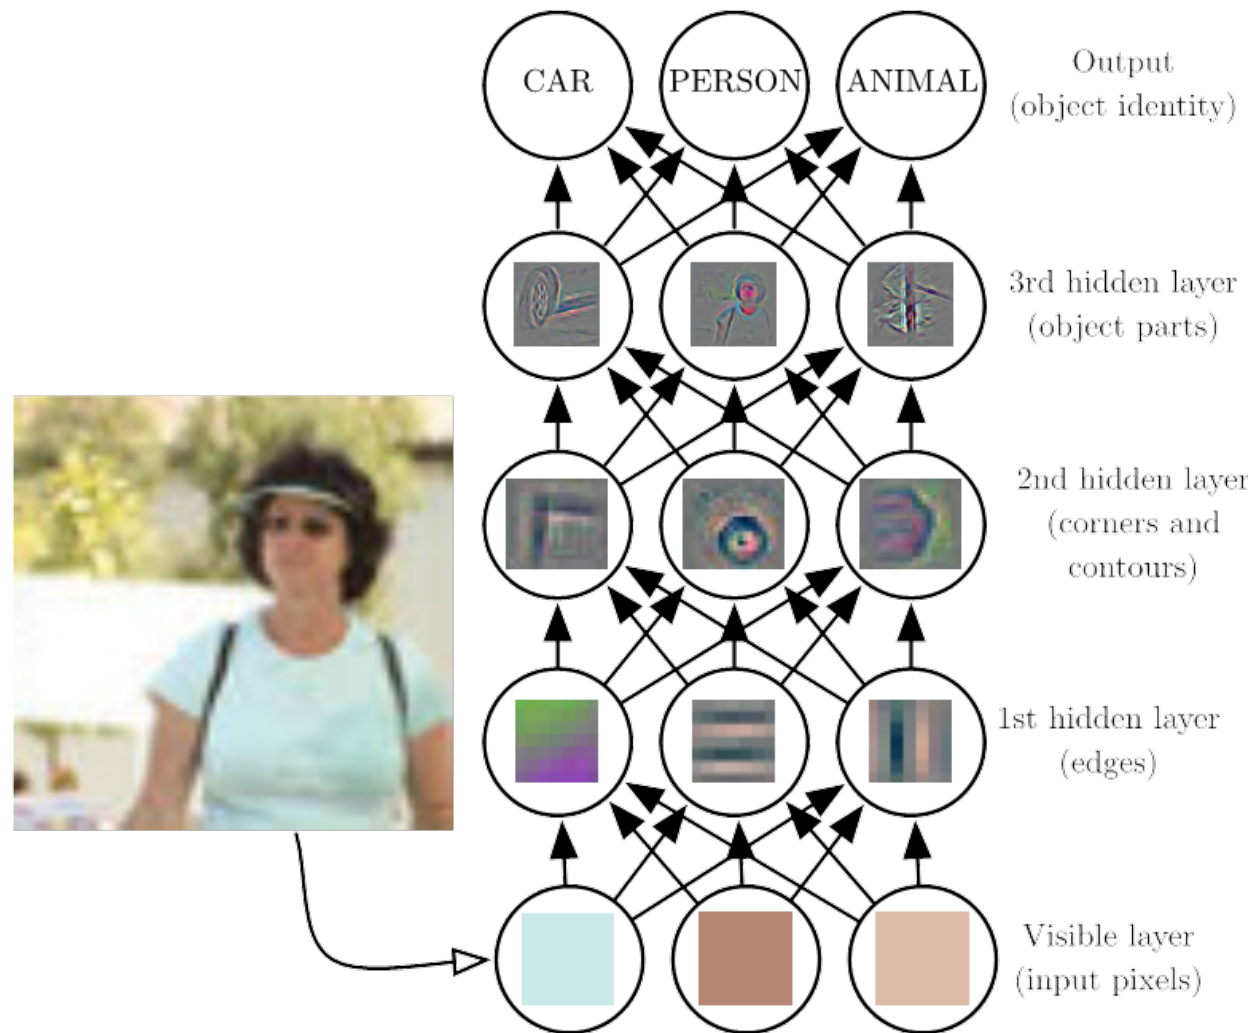
- Renaissance of artificial neural networks
 - Representation learning vs feature engineering
- Background
 - Linear Algebra, Optimization
 - Regularization
- Construction and training of layered learners
- Frameworks for deep learning

Representations matter

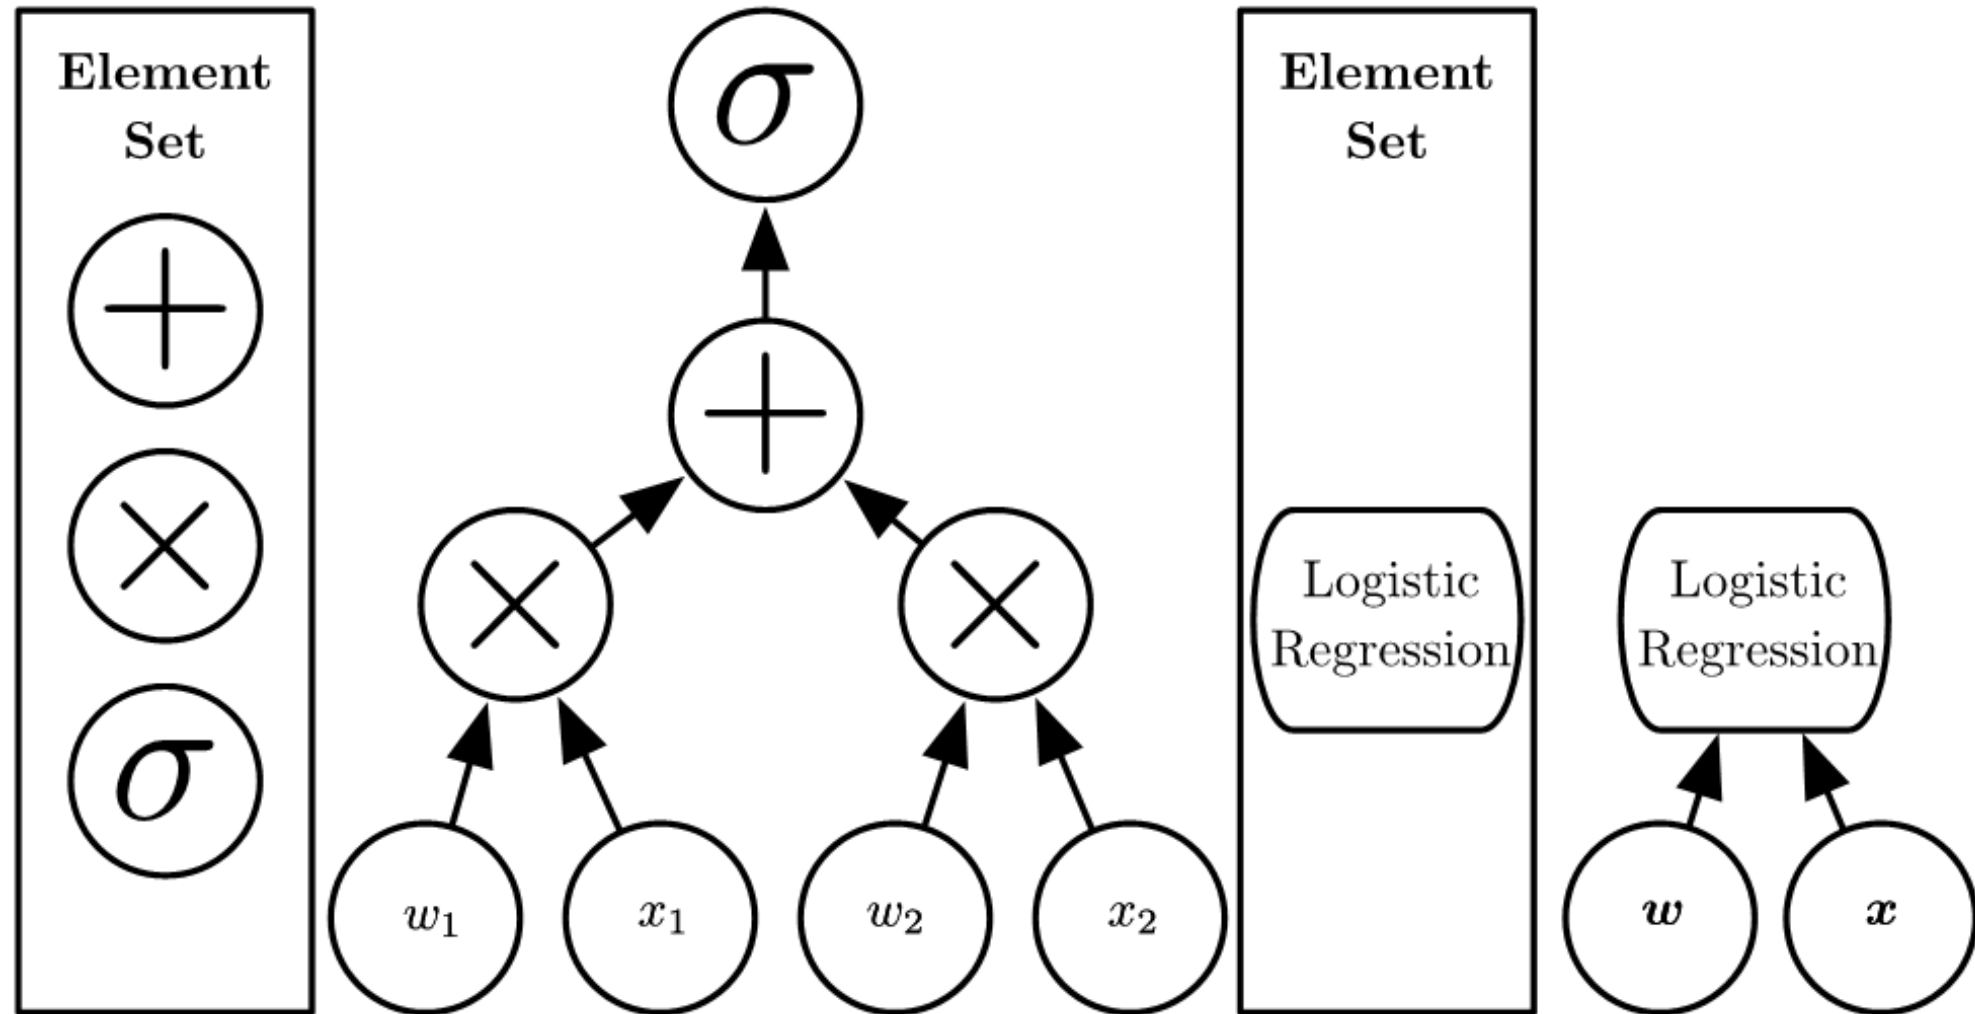


- Transform into the right representation
- Classify points simply by threshold on radius axis
- Single neuron with non-linearity can do this

Depth: layered composition

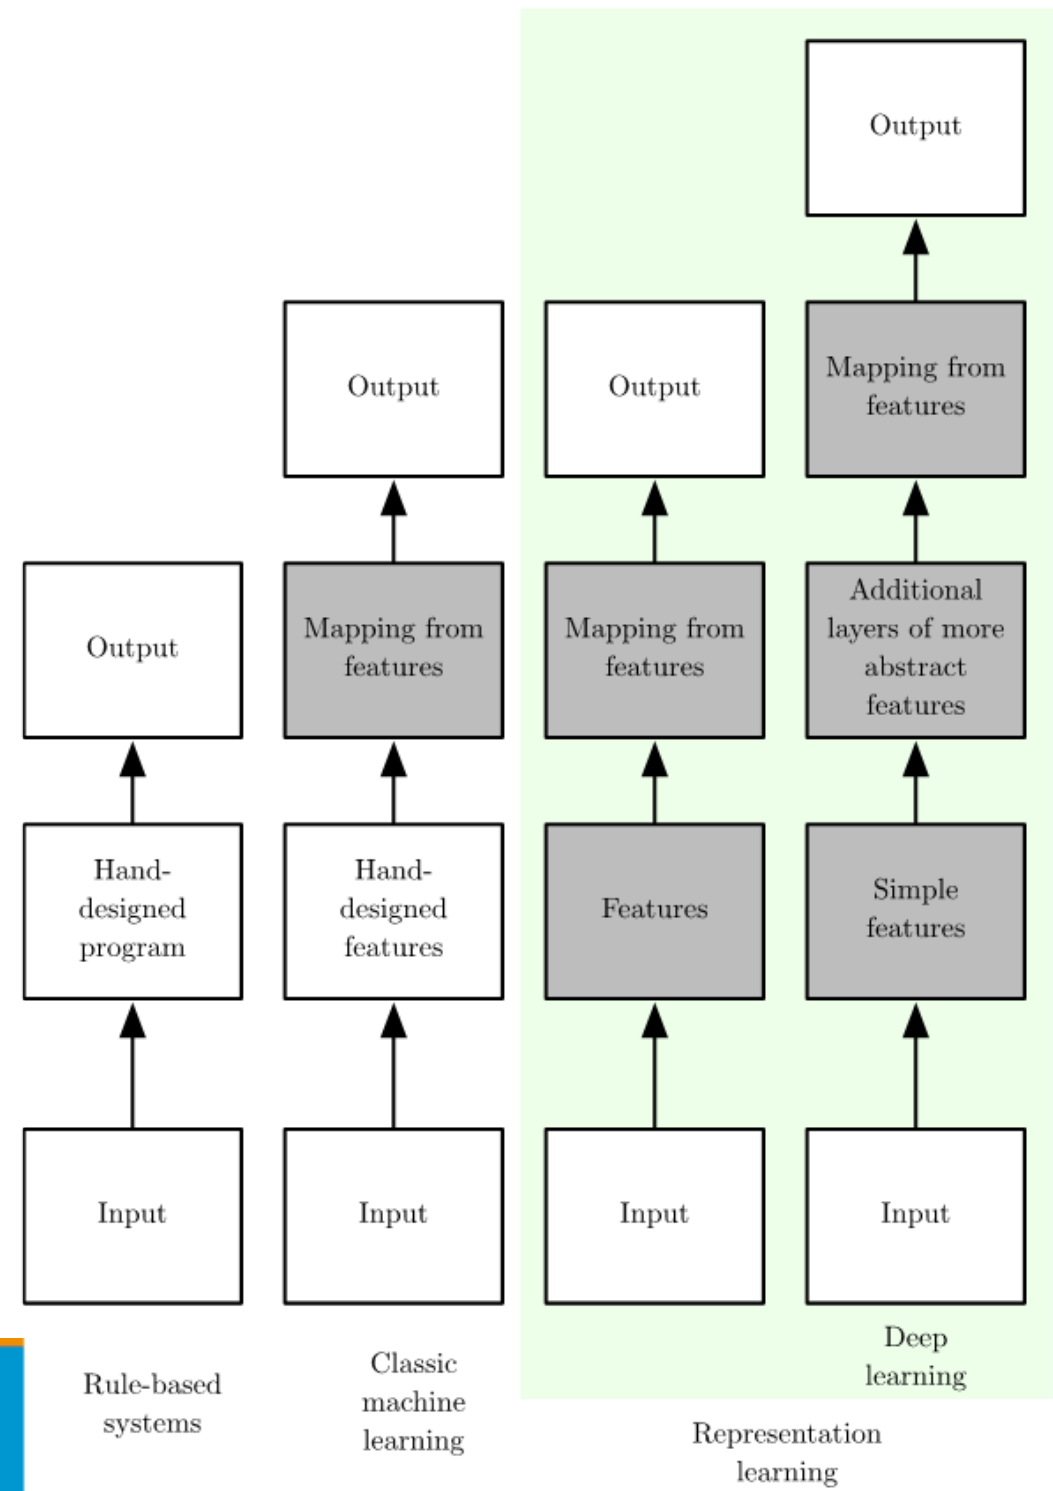


Computational graph




Components of learning

- Hand designed program
 - Input → Output
- Increasingly automated
 - Simple features
 - Abstract features
 - Mapping from features



Growing Dataset Size



MNIST dataset

8	9	0	1	2	3	4	7	8	9	0	1	2	3	4	5	6	7	8	6
4	2	6	4	7	5	5	4	7	8	9	2	9	3	9	3	8	2	0	5
0	1	0	4	2	6	5	3	5	3	8	0	0	3	4	1	5	3	0	8
3	0	6	2	7	1	1	8	1	7	1	3	8	9	7	6	7	4	1	6
7	5	1	7	1	9	8	0	6	9	4	9	9	3	7	1	9	2	2	5
3	7	8	2	3	4	5	6	7	8	9	0	1	2	3	4	5	6	7	0
1	2	3	4	5	6	7	8	9	8	1	0	5	5	1	9	0	4	1	9
3	8	4	7	7	8	5	0	6	5	5	3	3	3	9	8	1	4	0	6
1	0	0	6	2	1	1	3	2	8	8	7	8	4	6	0	2	0	3	6
8	7	1	5	9	9	3	2	4	9	4	6	5	3	2	5	5	9	4	1
6	5	0	1	2	3	4	5	6	7	8	9	0	1	2	3	4	5	6	7
8	9	0	1	2	3	4	5	6	7	8	9	6	4	2	6	4	7	5	5
4	7	8	9	2	9	3	9	3	8	2	0	9	8	0	5	6	0	1	0
4	2	6	5	5	5	4	3	4	1	5	3	0	8	3	0	6	2	7	1
1	8	1	7	1	3	8	5	4	2	0	9	7	6	7	4	1	6	8	4
7	5	1	2	6	7	1	9	8	0	6	9	4	9	9	6	2	3	7	1
9	2	2	5	3	7	8	0	1	2	3	4	5	6	7	8	0	1	2	3
4	5	6	7	8	0	1	2	3	4	5	6	7	8	9	2	1	2	1	3
9	9	8	5	3	7	0	7	7	5	7	9	9	4	7	0	3	4	1	4
4	7	5	8	1	4	8	4	1	8	6	6	4	6	3	5	7	2	5	9


Basics

Linear Algebra and Optimization


Linear Algebra

- Tensor is an array of numbers
 - Multi-dim: 0d scalar, 1d vector, 2d matrix/image, 3d RGB image

- Matrix (dot) product $C = AB$ $C_{i,j} = \sum_k A_{i,k} B_{k,j}$



- Dot product of vectors A and B
 - ($m = p = 1$ in above notation, $n=2$)




Linear algebra: Norms

- L^p norm

$$||\mathbf{x}||_p = \left(\sum_i |x_i|^p \right)^{\frac{1}{p}}$$

- Most popular norm: L2 norm, $p=2$
- L1 norm, $p=1$: $||\mathbf{x}||_1 = \sum_i |x_i|$.
- Max norm, infinite p : $||\mathbf{x}||_\infty = \max_i |x_i|$.





Nonlinearities

- ReLU


- Softplus

- Logistic Sigmoid




[Goodfellow, Bengio, Courville 2016]

Approximate Optimization




Gradient descent




Critical points



Minimum




Maximum



Saddle point



Saddle point – 1 and 2 derivative vanish



Poor conditioning:
1 deriv large in one and small in another direction

Tensorflow Playground


- <http://playground.tensorflow.org/>
 - Try out simple network configurations
- <https://cs.stanford.edu/people/karpathy/convnetjs/demo/classify2d.html>
 - Visualize linear and non-linear mappings

Regularization

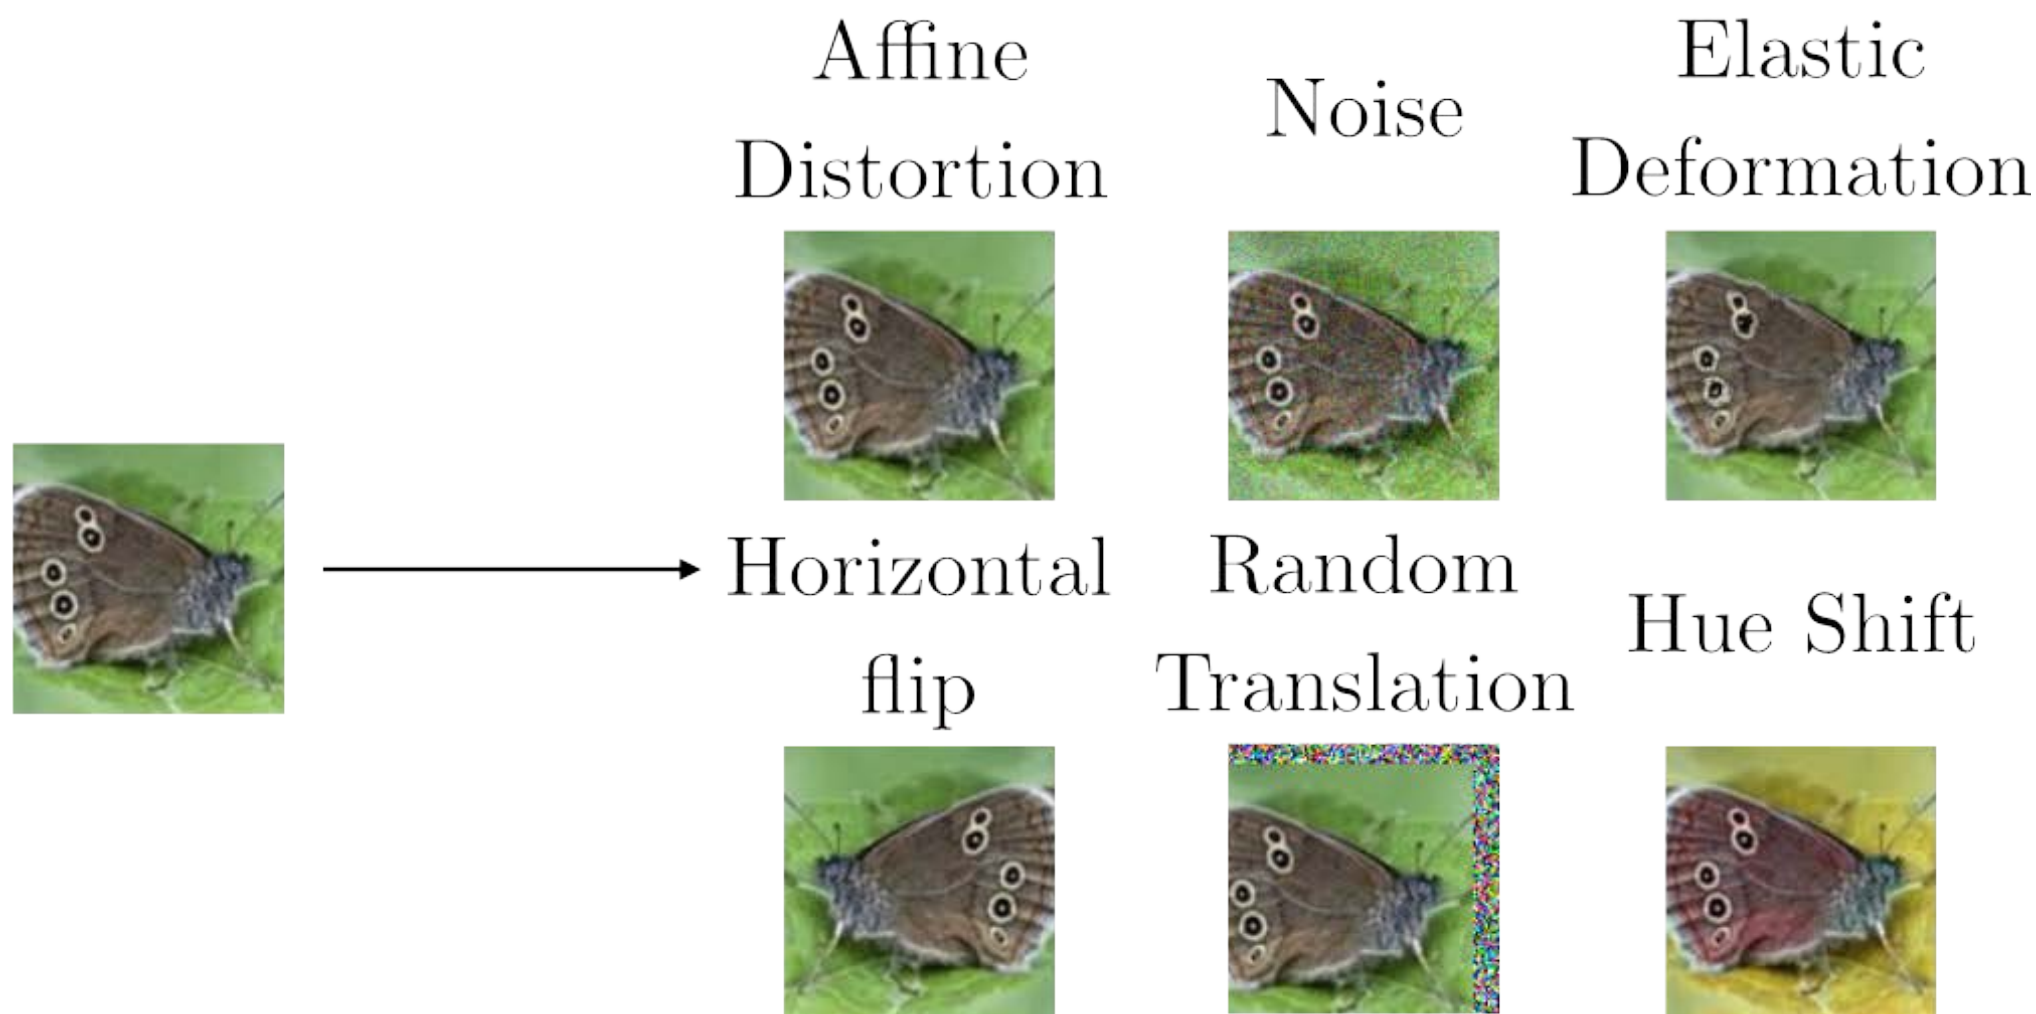
Lower generalization error
without impacting training error

Constrained optimization

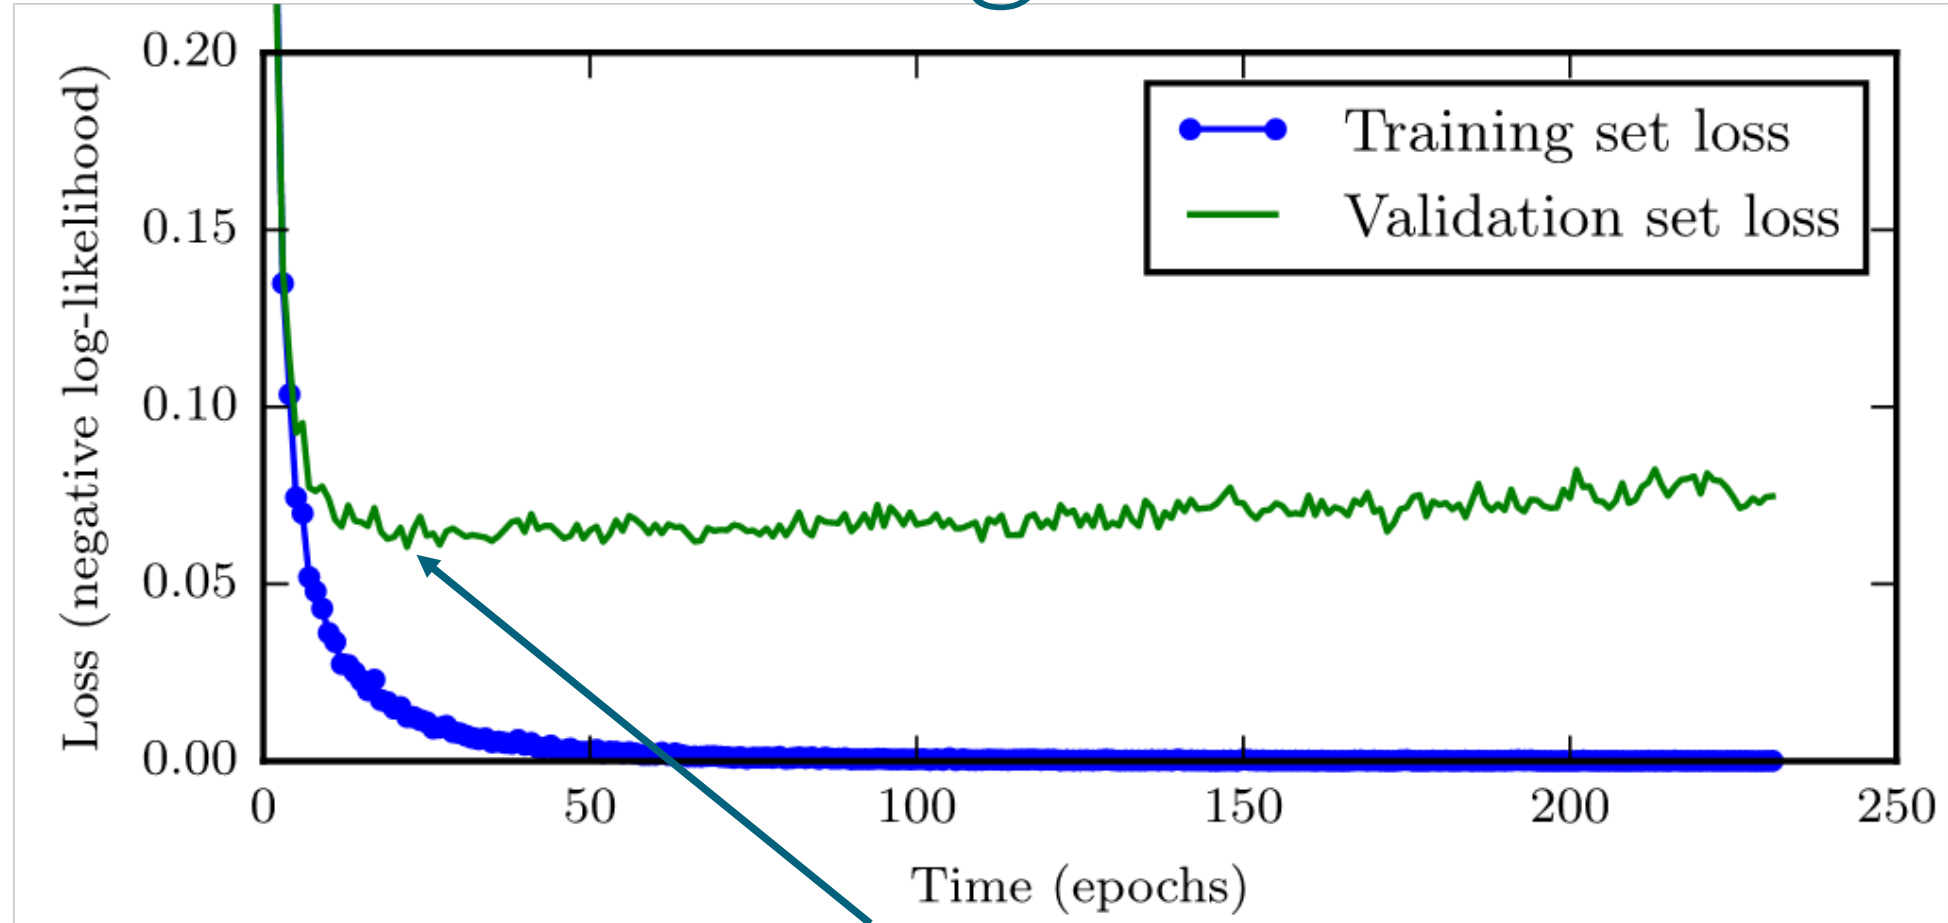
- Squared L2 encourages small weights
- L1 encourages sparsity of model parameters (weights)



Dataset augmentation



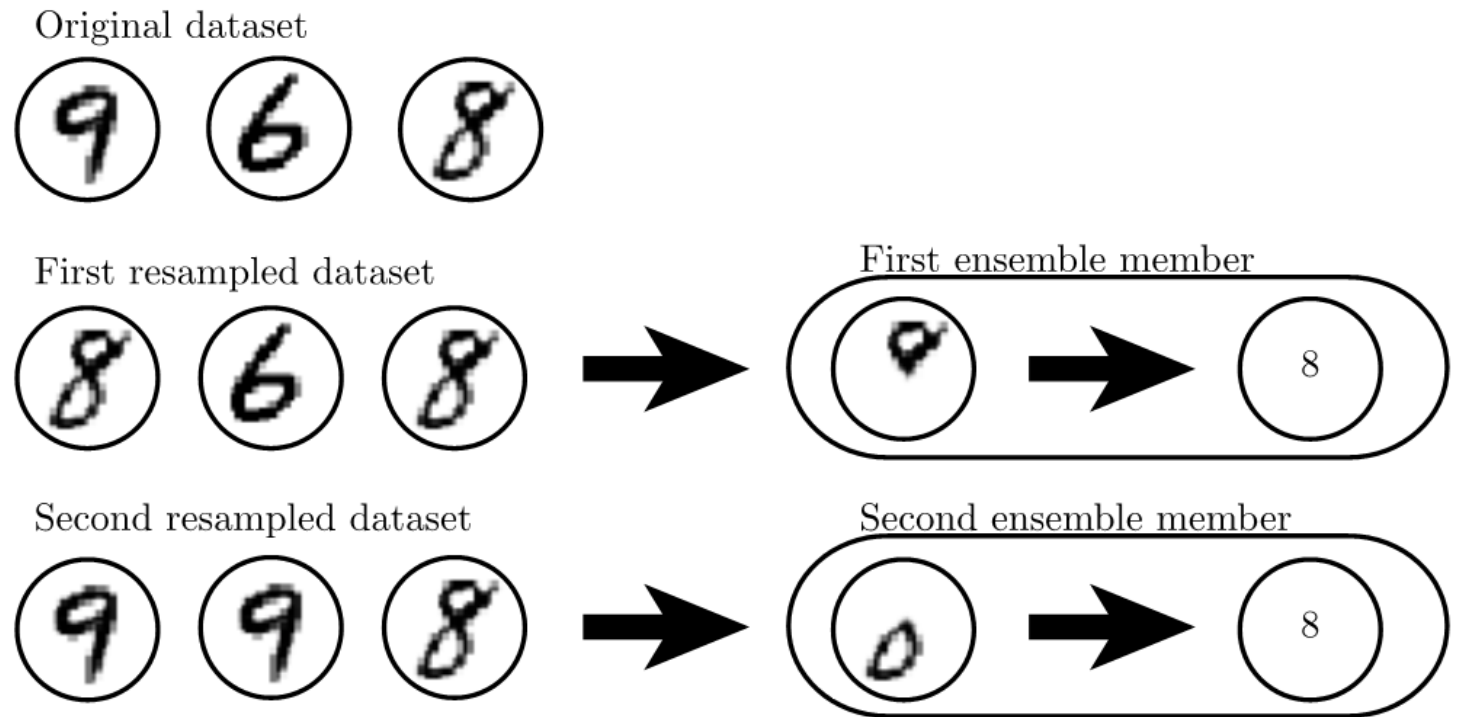
Learning curves



- Early stopping before validation error starts to increase



Bagging

- Average multiple models trained on subsets of the data
- First subset: learns top loop, Second subset: bottom loop




Dropout

- Random sample of connection weights is set to zero
- Train different network model each time
- Learn more robust, generalizable features



Multitask learning



- Shared parameters are trained with more data
- Improved generalization error due to increased statistical strength
- Missing components of y are masked from the loss function

Components of popular architectures


Convolution as edge detector



Input


1	-1
---	----

Kernel




Output


Gabor wavelets (kernels)



Directional second
derivative




Second derivative
(curvature)




Local average, first
derivative

Gabor-like learned kernels




- Features extractors provided by pretrained networks

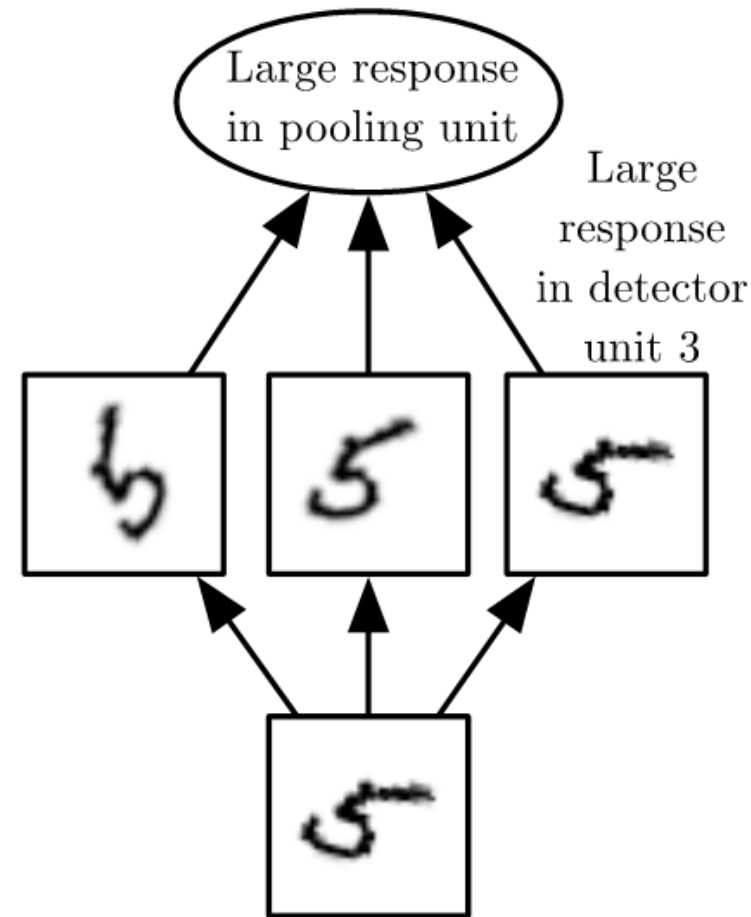
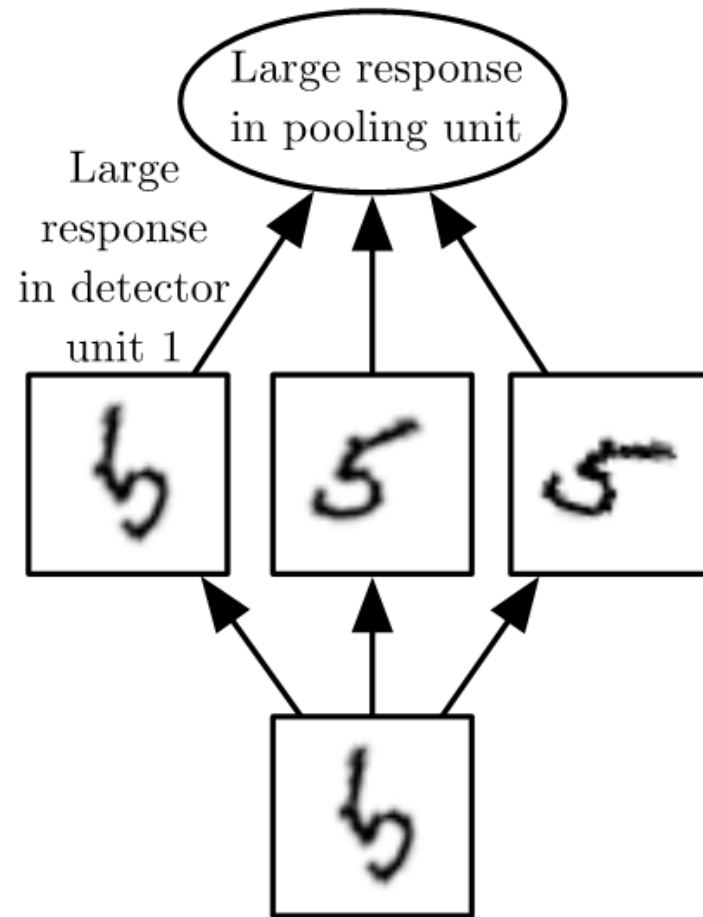
Max pooling translation invariance




- Take max of certain neighbourhood
- Often combined, followed by downsampling



Max pooling transform invariance



Types of connectivity



Local connection:
like convolution,
but no sharing

Choosing architecture family

- No structure → fully connected
- Spatial structure → convolutional
- Sequential structure → recurrent

Optimization Algorithm

- Lots of variants address choice of learning rate
- See [Visualization of Algorithms](#)
- AdaDelta and RMSprop often work well

Development strategy

- Identify needs: High accuracy or low accuracy?
- Choose metric
 - Accuracy (% of examples correct), Coverage (% examples processed)
 - Precision $TP/(TP+FP)$, Recall $TP/(TP+FN)$
 - Amount of error in case of regression
- Build end-to-end system
 - Start from baseline, e.g. initialize with pre-trained network
- Refine driven by data

Software for Deep Learning

Current Frameworks

- [Tensorflow / Keras](#)
- [PyTorch](#)
- DL4J
- Caffe (superseded by Caffe2, which is merged into PyTorch)
- [And many more](#)
- Most have CPU-only mode but much faster on NVIDIA GPU

Sources

- I. Goodfellow, Y. Bengio, A. Courville “Deep Learning” MIT Press 2016 [[link](#)]